

基于增量多任务深度学习的垂向分层水库藻类预测与藻华预警

A Multi-Task Deep Learning Framework with Incremental Prediction for Vertically Stratified Reservoir Algal Forecasting and Bloom Early Warning

Abstract

Reservoir algal blooms pose significant risks to drinking water safety, yet predicting bloom dynamics remains challenging due to strong temporal autocorrelation and vertical stratification. This study develops a multi-task deep learning framework (RAMS-Net) for a single stratified reservoir (2021–2025, 20 depth layers, 3-hour sampling). We demonstrate that (1) formulating the prediction target as *incremental change* ($\Delta = \text{conc}_{\{t+h\}} - \text{conc}_t$) rather than absolute concentration substantially improves skill over persistence baselines; (2) a shared GRU backbone with multi-task heads (forecasting, stratification, bloom warning) improves interval calibration; (3) daily-scale sampling with 7-day horizon better matches the physical process scale; and (4) probabilistic bloom warning provides $3.6\times$ longer lead time than threshold-based methods with half the false positives. The framework achieves CRPS skill of +22.1% over persistence with coverage of 0.766, and identifies 12 bloom events with median 19.5-day pre-event warning windows. These results provide a scientifically grounded and deployable approach for reservoir algae monitoring.

Keywords: reservoir algae; deep learning; incremental prediction; multi-task learning; bloom early warning; vertical stratification

1. Introduction

1.1 Background and Significance

Algal blooms in reservoirs threaten drinking water quality worldwide. Traditional monitoring relies on physical sampling and threshold-based warning, which are reactive and limited by vertical heterogeneity —

algae concentrations vary dramatically across depth layers (0.5–10 m), and surface measurements alone miss subsurface dynamics.

1.2 Related Work

Prior studies have applied deep learning to algal forecasting: LSTM-based prediction (Newanjiang Reservoir), temporal attention networks (GCN-TPA), and Transformer variants (Bloomformer-2). However, three gaps persist: (1) most models predict absolute concentration, which is dominated by strong autocorrelation and underperforms naive persistence; (2) few integrate multiple monitoring tasks (prediction, stratification, warning) in a shared architecture; (3) vertical profile data is underutilized.

1.3 Contributions

Incremental prediction formulation: predicting change (Δ) rather than absolute value outperforms persistence across all horizons, addressing the autocorrelation-dominated regime.

Multi-task shared backbone: GRU backbone with forecasting/stratification/warning heads improves both accuracy and interval calibration.

Daily-scale process alignment: matching sampling resolution to algal process timescale improves long-horizon forecasting.

Probabilistic bloom warning: calibrated probability output provides $3.6\times$ lead time over threshold methods with fewer false positives.

Vertical data utilization: systematic analysis of 20-layer profiles reveals stratification-state coupling during blooms.

2. Data and Methods

2.1 Study Site and Data

Single reservoir (site 0902, anonymized), 2021–03 to 2025–09. 258,542 algae measurements across 20 depth layers (0.5–10 m at 0.5 m intervals), 232,067 meteorological observations (10-min sampling). Key variables: water temperature, transmittance, algae concentrations (green/cyanobacteria/diatom/cryptophyte), total concentration, and six meteorological drivers (wind, temperature, humidity, pressure, rainfall).

Fig. 1 — Daily surface observations at the monitoring station (2021-03 to 2025-09; shaded: 12 N-defined bloom events)




Figure 1. Daily surface concentration and water temperature (z-scored) time series with the 12 identified bloom events shaded. Bloom events concentrate in 2021 (dominant) with secondary peaks in 2023.

Fig. 2 — Temporal distribution of the 12 N-defined bloom events




Figure 2. Gantt-style distribution of the 12 bloom events across the study period (2021–2025). 2021 events (N1–N4) highlighted in red, comprising 74% of total bloom days.

Fig. 3 — Vertical stratification of the water column




Figure 3. Weekly-mean z-scored heatmaps of 20-layer concentration (top) and water temperature (bottom). Concentration color scale clipped at $+3\sigma$ to highlight bloom layers.

2.2 Data Preprocessing

- Vertical timestamp alignment: depth layers have 4–5 min timestamp offsets (exact intersection = 0); floored to 3-hour grid (97.8% alignment).
- Meteorological alignment via `merge_asof` (nearest, 3-h tolerance).
- Daily aggregation (1D mean) for process-scale matching.

2.3 Bloom Event Definition

Bloom status = surface-band (0.5–3.0 m) median $>$ band-p90 AND ≥ 3 layers in 0.5–5.0 m $>$ their p90, sustained ≥ 2 days. This identifies 12 bloom events (2021 dominant: 74%), with 19.5-day median pre-event warning window.

2.4 Model Architecture

RAMS-Net: shared GRU backbone (hidden=64) $\rightarrow$ three task heads:

- M1 Forecasting: incremental target $\Delta = \text{conc}_{\{t+h\}} - \text{conc}_t$, 9-quantile output (q9)
- M2 Stratification: binary classification (stratified/not)
- M4 Bloom warning: bloom-status classification

Fig. 4 — RAMS-Net architecture (shared GRU backbone with M1/M2/M4 heads)




Figure 4. RAMS-Net multi-task architecture: shared GRU backbone with M1 incremental forecasting (q9 quantiles), M2 stratification, and M4 bloom warning heads. Multi-task loss with $w=(1, 3, 2)$.

2.5 Training

- **Two-stage:** Stage 1 single-task M1 (20 epochs) → Stage 2 freeze backbone, fine-tune multi-head (10 epochs)
- **Multi-task loss:** $L = w_1 \cdot L_q(M1) + w_2 \cdot CE(M2) + w_4 \cdot CE(M4)$, $w=(1, 3, 2)$
- **Daily scale:** T=30 lookback, H=7-day horizon
- **Rolling window evaluation:** 730d train / 90d test / 45d step (17 windows), 3 seeds

2.6 Evaluation

CRPS (continuous ranked probability score), coverage, p50 RMSE, relative skill vs persistence. Bloom warning: recall, lead time, false positives.

3. Results

3.1 Forecasting Performance




Figure 5. CRPS by forecast horizon comparing RAMS-Net (incremental, q9) with persistence, LightGBM-quantile, and XGBoost-quantile baselines. Error bars show ± 1 std across 3 seeds $\times$ 17 rolling windows.

Incremental prediction beats persistence across all horizons. RAMS-Net achieves CRPS skill of +22.1% over persistence (3-seed $\times$ 17-window), exceeding all statistical baselines: LightGBM quantile (+10.0%), XGBoost quantile (+4.6%). The only method with meaningful probabilistic calibration (coverage 0.766 vs 0.55 for GBDT baselines). All 7 horizons exceed +20% skill.

[Table 1: CRPS comparison table]

3.2 Interval Calibration

Fig. 6 — Predictive-interval coverage of quantile forecasts




Figure 6. Prediction interval coverage comparison: RAMS-Net formal (0.766) vs nominal 80% target; incremental quantiles (0.824) vs absolute-concentration quantiles (0.669).

Two-stage training with multi-task heads calibrates prediction intervals: coverage 0.766, approaching the 80% target. Absolute-concentration quantiles under-cover (0.669) due to seasonal level drift; incremental quantiles are better calibrated.

3.3 Bloom Warning

Fig. 7 — M4 bloom-warning threshold sensitivity (recall vs false alarms)




Figure 7. Bloom warning recall-false-positive tradeoff across probability threshold θ (0.3–0.8), with rule-based baseline reference point.

Fig. 8 — Warning lead time: probabilistic vs threshold-based warning




Figure 8. Median warning lead time: probabilistic warning ($\theta=0.5-0.7$) achieves 14-21 days versus 4 days for the rule-based threshold baseline — a $3.6\times$ improvement with half the false positives.

Probabilistic warning provides $3.6\times$ lead time. At threshold $\theta=0.5$, model recall 0.4 with median 14.5-day lead (13-16 days), 6 false-positive episodes. Threshold baseline (band p75): recall 1.0 but only 4-day lead, 14 false-positive episodes. Model trades half the false positives for $3.6\times$ longer lead.

3.4 Vertical Structure and Causality

Fig. 9 — M5 PCMCi+ causal structure (12-h grid, $\tau_{\max}=60$, $\alpha=0.05$)




Figure 9. PCMCi+ causal structure: algae concentration dominated by autoregression ($\tau=1$, $r=0.63$); meteorological drivers weak (wind_u $r \approx 0.07$); stratification decoupled during blooms.

PCMCi+ analysis shows algae concentration is dominated by its own autoregression ($\tau=1$ $r=0.63$); meteorological drivers have weak direct effects (wind_u $r \approx 0.07$); stratification decouples during blooms (surface-10m correlation drops to -0.07 vs 0.57 non-bloom).

3.5 Ablation Studies

Fig. 10 — Key ablation comparisons




Figure 10. Key ablations: two-stage training (ts_freeze best accuracy), incremental vs absolute target, multi-task vs single-task (calibration improvement).

- Incremental vs absolute target: incremental wins all horizons
- Multi-task vs single-task: multi-task improves calibration (0.67→0.81 coverage)
- Two-stage vs joint: two-stage ts_freeze achieves best accuracy
- Daily vs 3h scale: daily better for 7-day horizon

3.6 Framework Comparison

Fig. 11 — Framework comparison across 12 forecasting models




Figure 11. RMSE comparison across 12 frameworks: RAMS-Net (GRU multi-task, 3.64) optimal vs DLinear, TFT, XGBoost, LightGBM, Transformer, PatchTST.

RAMS-Net (GRU multi-task) outperforms all alternatives: DLinear, TFT, XGBoost, LightGBM, Transformer, PatchTST — GRU multi-task architecture is optimal for this data scale.

4. Discussion

4.1 Why Incremental Prediction Works

Algal concentration is a strong autoregressive process ($\tau=1$ $r=0.63$). Persistence (current value) already captures $\sim 80\%$ of the level. Modeling the *change* (Δ) lets the network learn corrections rather than redundantly modeling the level — this is the key insight, generalizable to other strongly autocorrelated environmental variables.

4.2 Multi-task Regularization

The stratification and warning heads act as regularizers, improving interval calibration without sacrificing point accuracy. This aligns with multi-task learning theory: auxiliary tasks provide inductive bias.

4.3 Process Scale Matching

Daily sampling with 7-day horizon matches algal growth (1-3 day doubling) and meteorological lag (13-30 days) timescales better than 3-hour sampling with 24-hour horizon. This suggests sampling resolution should be matched to the dominant process timescale.

4.4 Limitations

- **Information ceiling:** oracle probe shows the data approaches its predictable limit; pure autoregression captures most signal. Further improvement requires more data (multi-site, nutrients, inflow).
- **Direction predictability:** $P(\Delta > 0)$ is weakly discriminative (AUC 0.59); direction of change is largely unpredictable.
- **Seasonal drift:** interval coverage varies across windows (0.66–0.94) due to inter-annual non-stationarity; conformal calibration

only partially mitigates.

5. Conclusion

We developed RAMS-Net, a multi-task deep learning framework for vertically stratified reservoir algal forecasting. Key findings: (1) incremental prediction outperforms persistence by +22.1% CRPS skill; (2) daily-scale process matching improves long-horizon forecasting; (3) probabilistic bloom warning provides $3.6\times$ lead time with fewer false positives; (4) multi-task architecture calibrates prediction intervals. The framework offers a scientifically grounded, deployable approach for reservoir algae monitoring and early warning.

Acknowledgements

Data provided under confidentiality agreement; anonymized.

References

(To be completed with ≥ 15 references, 60% recent 5 years)

Rousso et al. Water Research 2020 — cyanobacteria forecasting review

Hipsey et al. GMD 2019 — General Lake Model

Carey et al. Inland Waters 2022 — FLARE forecasting

Thomas et al. WRR 2020 — iterative forecasting

Runge et al. — PCMCI+ causality

Romano et al. 2019 — conformal prediction (CQR)

...